




Audio Signal Processing for Machine Learning

Valerio Velardo


Problem




Problem



Problem



Problem



Applications

- Audio classification
- Speech recognition / speaker verification
- Audio denoising / audio upsampling
- Music Information Retrieval
 - Music Instrument Classification
 - Mood Classification
 - ...
- ...

Content

- Sound waves
- DAC / ADC
- Time- and frequency-domain audio features (e.g., rms, spectral centroid)
- Audio transformations
 - Fourier Transform / STFT
 - Constant-Q Transform
 - Mel Spectrograms
 - Chromograms
- ...

What should you expect?

- Theory
- Coding tutorials

Where do you I get the code/slides?




Technology stack



What you'll learn

- Get a deep understanding of audio data
- Familiarise with frequency/time-domain audio features
- Extract features from raw audio
- Recognise what audio features to use for ML applications
- Preprocess audio data for ML
- Understand (some!) math behind audio transformations
- Use *librosa* for your audio projects

Don't freak out!



Who's this series for?

- ML/DL engineers
- Computer science students
- Software engineers
- Music technologists
- Tech-oriented musicians

Prerequisites

- Intermediate Python programming

Join the community!



thesoundofai.slack.com